

Deterministic-classifier integration				10-shot exemplar tournament				Open-weight SOTA teacher			
0.8B	1.6 GB	67.6%	0%	500	10	10-shot	+1.58	31B	32 GB	\$0	0
Classifier params	CPU footprint	Test state acc	JSON fallback rate	Dialogues scored	Prompt levers	Winning template	Top-lever gain	Teacher params	Consumer GPU	Per-run API cost	Vendor lock-in
Qwen3.5-LoRA	no GPU for routing	4,304-turn split	from 21% at scale	10 cells × n=50	CoT, persona, N-best...	stage-balanced	length budget	Gemma 4 31B	RTX 5090	vs ~\$15 frontier	open weights

Headline outcome

2.14×	+29.45 pp	+19.25	+2.18	9.2×
State-acc multiplier	Absolute state-acc lift	Unified-score lead	Unified-score lead	Peak per-stage multiplier
vs paper baseline	25.94% to 55.39%	vs paper baseline	vs frontier ceiling	stage d (resolution)