

30×

faster routing per turn

Deterministic classifier consultant

A 0.8B-parameter Qwen3.5-LoRA on CPU replaces KELE's LLM consultant. ~50 ms forward pass vs. ~1.5 s API round-trip. Routing is deterministic and reproducible, not sampled.

500

dialogues across 10 levers

10-shot stage-balanced exemplar tournament

Swept length budget, persona, CoT scaffold, N-best rerank, +6 more. Winner: 10-shot stage-balanced exemplars, 2 per SocRule stage. Plain text generation removes KELE's JSON schema and its 21% fallback at scale.

\$0

per-run evaluation cost

Open-weight teacher on one consumer GPU

Gemma 4 31B served via vLLM on a single RTX 5090 (32 GB). No proprietary API, no rented hardware. Frontier API teachers cost ~\$15 per full n=681 evaluation.