

Within-architecture base-model ablation (n = 681)

STL's LoRA on top of GLM-4-9B-Chat-base · only variable = the adapter trained on SocratDataset

Surface-form gain is 1.68–1.85x
the pedagogical gain.
 Δ R-1 above base = contamination overlay.

