

Did SocratTeachLLM actually learn to teach better than its base model?

Within-architecture ablation (n = 681) · SocratTeachLLM = GLM-4-9B-Chat (base) + a LoRA adapter trained on KELE's SocratDataset

Result: LoRA adds +8 pp state acc
but +15 pp ROUGE-1 over its base.
The extra surface-form inflation
IS the memorization overlay.

