


Architecture: LLM consultant vs supervised classifier consultant

Replace the consultant axis; isolate routing from response generation

JSON schema · 21% Fallback rate at scale

Fallback = teacher emits malformed JSON; system reverts to a two-call recovery path



Deterministic SocRule lookup · No JSON, no sampling variance

Dimension	KELE	Ours
Per-run eval API cost	~\$15 / 1k turns	\$0 (CPU classifier)
Consultant footprint	API call	1.6 GB on CPU
Teacher footprint	9B · ~19 GB GPU	31B · 32 GB GPU
Routing determinism	LLM-sampled	Deterministic