


Architecture: LLM consultant vs. Supervised classifier

Replacing the consultant axis isolates routing from response generation



Dimension	KELE	Ours
Per-run eval API cost	~\$15 / 1k turns	\$0 (CPU classifier)
Consultant per-turn latency	~1.5 s (API round-trip)	~50 ms (CPU forward pass)
Teacher footprint	9B · ~19 GB GPU	31B · 32 GB GPU
Routing determinism	LLM-sampled	Deterministic